

PARTH RATHOD

CSP - 571 Fall 2021

A20458817

Recitation Answers

Chapter 4

Question 4)

a) Feature X is uniformly distributed on [0,1] and range of observation is within 10%. Hence available observations we can use for predictions is:

$$\begin{aligned}\text{For } X [0.55, 0.65] &= ((0.65 - 0.55) / (1-0)) * 100 \\ &= 10\%\end{aligned}$$

b) We are given with the set of observations with $P = 2$ features X_1 and X_2 that are also uniformly distributed over the same range such that $(X_1, X_2) \in [0,1] \times [0, 1]$

$$\begin{aligned}\text{So, the fraction of available observations that we will use to make} \\ \text{prediction is given by:} \\ &= (10\%) * (10\%) \\ &= 1\%\end{aligned}$$

c) We are given with the set of observations with $P = 100$ features and all of which are uniformly distributed over the range between 0 and 1.

$$\begin{aligned}\text{So, the fraction of available observations that we will use to make} \\ \text{prediction is given by:} \\ &= (0.1)^{100} * 100 \\ &= (10)^{-98} \%\end{aligned}$$

d) From the above observations we can conclude that as the number of features increases, the percentage of observations that are used to predict KNN becomes very small. Hence, more features leads to fewer neighbors.

e)

$$\begin{aligned}\text{For } P = 1 &\Rightarrow \text{length} = (0.10) \\ \text{For } P = 2 &\Rightarrow \text{length} = (0.10)^{1/2} = 0.316 \\ \text{For } P = 100 &\Rightarrow \text{length} = (0.10)^{1/100} = 0.977\end{aligned}$$

From the above, we can conclude that when we wish to make predictions for the test observations that contains on average 10% of the training observations, then if we have large number of features, it will be better to include all the features.

Question 6)

a)

Logistic Regression with multiple variables is given by:-

$$p(X) = e^{(\beta_0 + X_1\beta_1 + X_2\beta_2)} / 1 + e^{(\beta_0 + X_1\beta_1 + X_2\beta_2)}$$

$$\beta_0 = -6, \beta_1 = 0.05, \beta_2 = 1, e = 2.71828$$

$$X_1 = 40, X_2 = 3.5$$

Substituting in Equation gives

$$p(X) = 37.755$$

b)

Logistic Regression with multiple variables is given by:-

$$p(X) = e^{(\beta_0 + X_1\beta_1 + X_2\beta_2)} / 1 + e^{(\beta_0 + X_1\beta_1 + X_2\beta_2)}$$

$$\beta_0 = -6, \beta_1 = 0.05, \beta_2 = 1, e = 2.71828$$

$$P(x) = 0.5, X_2 = 3.5$$

Substituting in equation gives

$$X_1 = 50\text{hrs}$$

Question 7)

9)

7) Bayes Theorem:-

$$P_k(x) = \frac{\pi_k f_k(x)}{\sum_{i=1}^K \pi_i f_i(x)}$$

$$\text{where } f_k(x) = \frac{1}{\sqrt{2\pi}\sigma_k} e^{\left(\frac{-1}{2\sigma_k^2}(x - \mu_k)^2\right)}$$

$$\therefore \pi_{\text{yes}} = 0.8, \pi_{\text{no}} = 0.2, \mu_{\text{yes}} = 10, \mu_{\text{no}} = 0, \sigma^2 = 36$$

Substituting in equation.
we get,

$$\underline{P_{\text{yes}}(x) = 0.752}$$

Question 9)

9]

a)

$$\text{odds} = \frac{P(x)}{1 - P(x)}$$

$$\therefore 0.37 = \frac{P(x)}{1 - P(x)}$$

$$\therefore P(x) + 0.37 P(x) = 0.37$$

$$\therefore P(x) = 0.27$$

$$\therefore \underline{\underline{P(x) = 27\%}}$$

b) $P(x) = 0.16$

$$\therefore 1 - P(x) = 1 - 0.16 = 0.84$$

$$\therefore \text{odds} = \frac{0.16}{0.84}$$

$$\therefore \underline{\underline{\text{Odds} = 0.19}}$$

Chapter 5

Question 2)

a)

Given: total number of observations = n
 Since bootstrap allows sampling with replacement
 Every observation in original sample is independent and has equal probability to appear in each bootstrap observation
 Probability that first bootstrap observation is the j th observation from original sample is = $1/n$
 Probability that first bootstrap observation is not the j th observation from original sample is = $1 - (1/n)$

b)

Given: total number of observations = n
 Since bootstrap allows sampling with replacement
 Every observation in original sample has equal probability to appear in each bootstrap observation
 Probability that Second bootstrap observation is the j th observation from original sample is = $1/n$
 Probability that Second bootstrap observation is not the j th observation from original sample is = $1 - (1/n)$

c)

$p(\text{jth observation is not in the first bootstrap sample}) = 1 - (1/n)$
 $p(\text{jth observation is not in the second bootstrap sample}) = 1 - (1/n)$
 \vdots
 $p(\text{jth observation is not in the nth bootstrap sample}) = 1 - (1/n)$

 $P(\text{jth observation is not in the bootstrap sample}) = (1 - (1/n)) * (1 - (1/n)) * \text{upto } n \text{ times}$

$$= (1 - (1/n))^n$$

d)

From above answer we can replace n by 5 and get the answer
 $P(\text{jth observation is in the bootstrap sample}) = (1 - (1/5))^5$

$$= 0.67$$

e)

$P(\text{jth observation is in the bootstrap sample}) = 1 - (1 - (1/n))^n$


$$= 1 - (1 - (1/100))^{100}$$

$$= 0.634$$

f)

$$\begin{aligned}
 P(\text{jth observation is in the bootstrap sample}) &= 1 - (1 - (1/n))^n \\
 &= 1 - (1 - (1/1000))^1000 \\
 &= 0.632
 \end{aligned}$$

g) `plot(1:100000, 1-(1-1/1:100000)^(1:100000), xlab= "Value of n",ylab = "Probabilities")`



Probability quickly reaches to 0.62 and remains constant as n tends to infinity

h)

```

data <- rep(NA, 10000)
for (i in 1:10000)
{
  data[i] <- sum(sample(1:100, rep = TRUE) == 4) > 0
}
mean(data)

```

Answer:- 0.6289

Question 3)

a)

In k-fold cross validation, set of observations is randomly divided into k groups of approx equal size. The first fold is treated as validation set and remaining k-1 as where the fit method is used. The MSE is then computed. This process is repeated k times, each time a different group of observation is treated as validation set. The k-fold CV is computed by doing average of these values, $CV(k) = \frac{1}{k} * \text{Summation}(MSE(i) \text{ for } i = 1 \text{ to } k)$

b)

Validation set Approach:

Advantages:

This approach is conceptually simple.

This approach is easy to implement.

Disadvantages:

Validation estimate of test error rate can be highly variable, depends on which observations are included in training set and which observations are included in validation set.

In this approach, only observations which are included in training set are used to fit the model and model can perform worse if fewer.

LOOCV:

Advantages:

The LOOCV cross-validation method is a subset of k-fold cross-validation with $k=n$.

In this technique, each observation is only considered once in a validation set and $(n-1)$ times in a training set.

This approach gives approximately unbiased estimates of the test error since each training set contain $(n-1)$ observations.

Disadvantages:

It necessitates fitting the potentially computationally costly model n times, as opposed to k-fold cross-validation, which necessitates fitting the model just k times, where $k < n$.

This method has a greater variance than k-fold cross-validation (since we are averaging the outputs of n fitted models trained on a virtually similar set of observations, these outputs are highly correlated, and the mean of these outputs is higher).

#Problem 1

```
library(readr)
library(data.table)
library(corrplot)
```

```
## corrplot 0.90 loaded
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(ROCR)
```

```
abaloneURL = "https://archive.ics.uci.edu/ml/machine-learning-databases/abalone/abalone.data"
abaloneData = fread(abaloneURL, header = FALSE)
abaloneHeader = c("Sex", "Length", "Diameter", "Height", "Whole weight", "Shucked weight", "Viscera weight", "Shell weight", "Rings")
colnames(abaloneData) = abaloneHeader
#Display Summary:
summary(abaloneData)
```

```
##      Sex      Length      Diameter      Height
## Length:4177   Min.    :0.075   Min.    :0.0550   Min.    :0.0000
## Class :character 1st Qu.:0.450   1st Qu.:0.3500   1st Qu.:0.1150
## Mode  :character Median :0.545   Median :0.4250   Median :0.1400
##              Mean  :0.524   Mean  :0.4079   Mean   :0.1395
##              3rd Qu.:0.615   3rd Qu.:0.4800   3rd Qu.:0.1650
##              Max.   :0.815   Max.   :0.6500   Max.    :1.1300
## Whole weight  Shucked weight  Viscera weight  Shell weight
## Min.    :0.0020   Min.    :0.0010   Min.    :0.0005   Min.    :0.0015
## 1st Qu.:0.4415   1st Qu.:0.1860   1st Qu.:0.0935   1st Qu.:0.1300
## Median :0.7995   Median :0.3360   Median :0.1710   Median :0.2340
## Mean   :0.8287   Mean   :0.3594   Mean   :0.1806   Mean   :0.2388
## 3rd Qu.:1.1530   3rd Qu.:0.5020   3rd Qu.:0.2530   3rd Qu.:0.3290
## Max.    :2.8255   Max.    :1.4880   Max.    :0.7600   Max.    :1.0050
## Rings
## Min.    : 1.000
## 1st Qu.: 8.000
## Median : 9.000
## Mean    : 9.934
## 3rd Qu.:11.000
## Max.    :29.000
```



```
#Returns the indices when the following condition is true in which()
condition = which(abaloneData$Sex!="I")
#All the Rows with Sex = I are removed
abaloneData2 = abaloneData[condition]
#check for all sex categories in new dataset:
unique(abaloneData2$Sex)
```

```
## [1] "M" "F"
```

```
#As the response variable need to be numeric, converting categorical Sex to a factor
abaloneData2$Sex = factor(abaloneData2$Sex)
#check for the change:
str(abaloneData2)
```

```
## Classes 'data.table' and 'data.frame':  2835 obs. of  9 variables:
## $ Sex          : Factor w/ 2 levels "F","M": 2 2 1 2 1 1 2 1 1 2 ...
## $ Length       : num  0.455 0.35 0.53 0.44 0.53 0.545 0.475 0.55 0.525 0.43 ...
## $ Diameter     : num  0.365 0.265 0.42 0.365 0.415 0.425 0.37 0.44 0.38 0.35 ...
## $ Height       : num  0.095 0.09 0.135 0.125 0.15 0.125 0.125 0.15 0.14 0.11 ...
## $ Whole weight : num  0.514 0.226 0.677 0.516 0.777 ...
## $ Shucked weight: num  0.2245 0.0995 0.2565 0.2155 0.237 ...
## $ Viscera weight: num  0.101 0.0485 0.1415 0.114 0.1415 ...
## $ Shell weight  : num  0.15 0.07 0.21 0.155 0.33 0.26 0.165 0.32 0.21 0.135 ...
## $ Rings        : int   15 7 9 10 20 16 9 19 14 10 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
#Creating 80-20 Training Testing Split, createDataPartition() returns the indices
trainIndex = createDataPartition(y = abaloneData2$Sex, p = 0.8, list = FALSE)
#Training data
trainData = abaloneData2[trainIndex,]
#Testing data (note the minus sign)
testData = abaloneData2[-trainIndex,]
```

```
#Predicting Sex using glm
model <- glm(Sex~.,family=binomial,data=trainData)
#Summary of our model
summary(model)
```

```
##
## Call:
## glm(formula = Sex ~ ., family = binomial, data = trainData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8422  -1.2014   0.8803   1.1195   1.4946
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    2.746856   0.517707   5.306 1.12e-07 ***
## Length        -1.404160   2.281488  -0.615  0.53825
## Diameter       -5.595478   2.738485  -2.043  0.04103 *
## Height        -1.457616   2.549472  -0.572  0.56750
## `Whole weight`  0.239246   0.845490   0.283  0.77720
## `Shucked weight` 2.877994   1.007477   2.857  0.00428 **
## `Viscera weight` -2.167277   1.420058  -1.526  0.12696
## `Shell weight` -0.230110   1.320678  -0.174  0.86168
## Rings          -0.001951   0.017841  -0.109  0.91290
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 3131.7  on 2268  degrees of freedom
## Residual deviance: 3072.8  on 2260  degrees of freedom
## AIC: 3090.8
##
## Number of Fisher Scoring iterations: 4
```

```
#Coefficients of our model
coef(model)
```

```
##      (Intercept)      Length      Diameter      Height
##      2.746855671    -1.404159621    -5.595477554    -1.457616395
## `Whole weight` `Shucked weight` `Viscera weight` `Shell weight`
##      0.239245893      2.877993519      -2.167277467      -0.230110287
##              Rings
##     -0.001951414
```

```
#Confidence Interval
confint(model)
```

```
## Waiting for profiling to be done...
```


##	2.5 %	97.5 %
## (Intercept)	1.74607346	3.77677275
## Length	-5.87792848	3.07173332
## Diameter	-10.98329493	-0.23982555
## Height	-6.46211360	3.60035582
## `Whole weight`	-1.41975707	1.90635573
## `Shucked weight`	0.90691967	4.86416179
## `Viscera weight`	-4.96292146	0.61144730
## `Shell weight`	-2.83273830	2.35616736
## Rings	-0.03694459	0.03303912

From the above observations for confidence interval, we can see that, all the predictors, except the “Shucked Weight” contain 0 within their confidence interval range which is in line with the basic assumption of Null Hypothesis which says, No relationship between X and Y. The only attribute with a substantially low p-value is “Shucked Weight” with a p-value of 0.00105 and all the other attributes have a high p-value. Thus, we can conclude that, there is no relationship between the predictors and the response variable (except the “Shucked Weight”) and Null Hypothesis holds true for all the predictors except the “Shucked Weight”. Hence Shucked Weight is the only predictor which is relevant.

```
#Predict [By setting the parameter type='response', R will output probabilities in the form of P (y=1|X)]
probs = predict(model, testData, type = "response")
#Using a 50% cut-off factor i.e probabilities > 0.5 are Males and rest are Females
resultSet = ifelse(probs > 0.5,"M","F")
resultSet2 = factor(resultSet)
#Creating a confusion matrix
confusionMatrix(resultSet2, testData$Sex)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  F    M
##           F 101  93
##           M 160 212
##
##           Accuracy : 0.553
##           95% CI : (0.511, 0.5945)
##           No Information Rate : 0.5389
##           P-Value [Acc > NIR] : 0.2638
##
##           Kappa : 0.0836
##
##           McNemar's Test P-Value : 3.334e-05
##
##           Sensitivity : 0.3870
##           Specificity : 0.6951
##           Pos Pred Value : 0.5206
##           Neg Pred Value : 0.5699
##           Prevalence : 0.4611
##           Detection Rate : 0.1784
##           Detection Prevalence : 0.3428
##           Balanced Accuracy : 0.5410
##
##           'Positive' Class : F
##
```

```
#Plotting the ROC Curve
roc.pred = prediction(probs,testData$Sex)
roc.perf = performance(roc.pred, measure = "tpr", x.measure = "fpr")
plot(roc.perf)
abline(0,1)
```



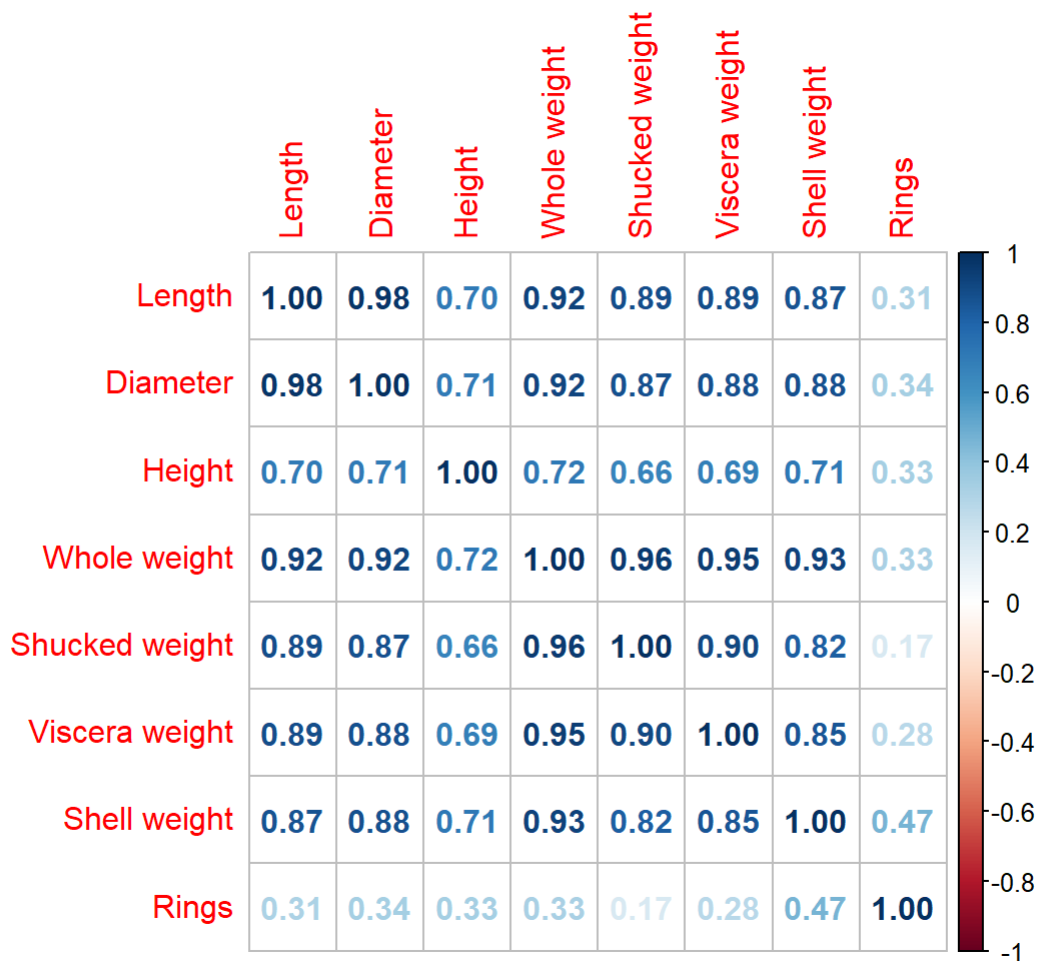
```
auc.perf = performance(roc.pred, measure = "auc")  
cat("Area Under the Curve: ")
```

```
## Area Under the Curve:
```

```
auc.perf@y.values
```

```
## [[1]]  
## [1] 0.5799259
```

```
#Plotting the correlations between the predictors  
cm = cor(abaloneData2[,-1])  
corrplot(cm, method = "number")
```



The above figure clearly shows that there is a positive linear connection for all factors. Only the Rings predictor has a weak uphill (positive) association, whereas the others have a high uphill (positive) relationship.

#Problem 2

```
library(data.table) #Data Import
library(e1071)      #Naïve Bayes
```

```
#Setting up the URL for data import
mushroomURL = "https://archive.ics.uci.edu/ml/machine-learning-databases/mushroom/agaricus-lepiota.data"
mushroomData = fread(mushroomURL,header=FALSE)
#Adding the headers
mushroomHeader = c("Class","cap-shape","cap-surface","cap-color","bruises","odor","gill-attachment","gill-spacing","gill-size","gill-color","stalk-shape","stalk-root","stalk-surface-above-ring","stalk-surface-below-ring","stalk-color-above-ring","stalk-color-below-ring","veil-type","veil-color","ring-number","ring-type","spore-print-color","population","habitat")
colnames(mushroomData) = mushroomHeader
#Display Summary:
summary(mushroomData)
```

```
##      Class      cap-shape      cap-surface      cap-color
## Length:8124      Length:8124      Length:8124      Length:8124
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## bruises          odor          gill-attachment gill-spacing
## Length:8124      Length:8124      Length:8124      Length:8124
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## gill-size        gill-color      stalk-shape      stalk-root
## Length:8124      Length:8124      Length:8124      Length:8124
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## stalk-surface-above-ring stalk-surface-below-ring stalk-color-above-ring
## Length:8124      Length:8124      Length:8124
## Class :character Class :character Class :character
## Mode :character Mode :character Mode :character
## stalk-color-below-ring veil-type      veil-color
## Length:8124      Length:8124      Length:8124
## Class :character Class :character Class :character
## Mode :character Mode :character Mode :character
## ring-number      ring-type      spore-print-color population
## Length:8124      Length:8124      Length:8124      Length:8124
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## habitat
## Length:8124
## Class :character
## Mode :character
```

```
#Class Distribution
table(mushroomData$Class)
```

```
##
##      e      p
## 4208 3916
```

```
#Converting Class Attribute to a factor
mushroomData$Class = factor(mushroomData$Class)
#Dimensions
dim(mushroomData)
```

```
## [1] 8124 23
```

```
#Structure
str(mushroomData)
```

```
## Classes 'data.table' and 'data.frame': 8124 obs. of 23 variables:
## $ Class : Factor w/ 2 levels "e","p": 2 1 1 2 1 1 1 1 2 1 ...
## $ cap-shape : chr "x" "x" "b" "x" ...
## $ cap-surface : chr "s" "s" "s" "y" ...
## $ cap-color : chr "n" "y" "w" "w" ...
## $ bruises : chr "t" "t" "t" "t" ...
## $ odor : chr "p" "a" "l" "p" ...
## $ gill-attachment : chr "f" "f" "f" "f" ...
## $ gill-spacing : chr "c" "c" "c" "c" ...
## $ gill-size : chr "n" "b" "b" "n" ...
## $ gill-color : chr "k" "k" "n" "n" ...
## $ stalk-shape : chr "e" "e" "e" "e" ...
## $ stalk-root : chr "e" "c" "c" "e" ...
## $ stalk-surface-above-ring: chr "s" "s" "s" "s" ...
## $ stalk-surface-below-ring: chr "s" "s" "s" "s" ...
## $ stalk-color-above-ring : chr "w" "w" "w" "w" ...
## $ stalk-color-below-ring : chr "w" "w" "w" "w" ...
## $ veil-type : chr "p" "p" "p" "p" ...
## $ veil-color : chr "w" "w" "w" "w" ...
## $ ring-number : chr "o" "o" "o" "o" ...
## $ ring-type : chr "p" "p" "p" "p" ...
## $ spore-print-color : chr "k" "n" "n" "k" ...
## $ population : chr "s" "n" "n" "s" ...
## $ habitat : chr "u" "g" "m" "u" ...
## - attr(*, ".internal.selfref")=<externalptr>
```

#Finding the number of missing values

```
cat("Number of missing values = ",sum(mushroomData=="?"))
```

```
## Number of missing values = 2480
```

#New dataset with removed missing values

```
mushroomData2 = mushroomData[mushroomData$`stalk-root`!="?"]
```

There is no damage in removing the observations with missing values because we have plenty after omitting the ones with missing values.

#Creating a split

```
trainSize = floor(0.80*nrow(mushroomData2))
trainIndex = sample(nrow(mushroomData2), size = trainSize)
trainData = mushroomData2[trainIndex,]
testData = mushroomData2[-trainIndex,]
```



```

#Naive Bayes
model = naiveBayes(trainData[,-1],trainData$Class)
#Prediction on Testing Data
testPred = predict(model,testData[,-1])
#Prediction on Training Data
trainPred = predict(model,trainData[,-1])
#Accuracy of Testing Model
cat("Accuracy of Testing Model: ",mean(testPred == testData$Class)*100,"%")

```

```
## Accuracy of Testing Model: 95.39415 %
```

```

#Accuracy of Training Model
cat("Accuracy of Training Model: ",mean(trainPred == trainData$Class)*100,"%")

```

```
## Accuracy of Training Model: 95.43743 %
```

```

#Confusion Matrix
table(testPred, testData$Class)

```

```

##
## testPred  e   p
##           e 691 47
##           p  5 386

```

From the above the confusion matrix we have the following: TP = 691, FP = 47, FN = 5 and TN = 386

#Problem 3

```

library(readr)
library(data.table)
library(caret)

```

```

hd_URL = "http://archive.ics.uci.edu/ml/machine-learning-databases/00243/yacht_hydrodynamics.dat
a"
hd_data = fread(hd_URL, header = FALSE)

hd_Header = c("longitudinalPos","prismaticCoef","LDR","BDR","LBR","froudeNo","Residuary")
colnames(hd_data) = hd_Header

#Direct use of attribute names in code
attach(hd_data)

```

```
#Creating 80-20 Training Testing Split, createDataPartition() returns the indices  
trainIndex = createDataPartition(y = hd_data$Residuary , p = 0.8, list = FALSE)
```

```
#Training data  
trainData = hd_data[trainIndex,]
```

```
#Testing data (note the minus sign)  
testData = hd_data[-trainIndex,]
```

```
#Training fit for linear model  
linearModel1 = lm(hd_data$Residuary~hd_data$longitudinalPos + hd_data$prismaticCoef + hd_data$LD  
R + hd_data$BDR + hd_data$BDR + hd_data$LBR +hd_data$froudeNo, data = trainData)
```

```
# Function to compute MSE  
MSE = function(yActual, yPred)  
{  
  return (mean((yActual - yPred)^2))  
}
```

```
mse1 = MSE(hd_data$Residuary, linearModel1$fitted.values )
```

```
# Summarize the results  
cat("Training MSE: ", mse1)
```

```
## Training MSE: 78.45015
```

```
cat("Training RMSE: ", sqrt(mse1))
```

```
## Training RMSE: 8.857209
```

```
cat("Training R-squared: ",summary(linearModel1)$r.sq)
```

```
## Training R-squared: 0.6575638
```

```
# Define training control  
train.control = trainControl(method = "boot", number = 1000)  
  
# Train the model  
linearModel2 = train(Residuary~., data = trainData, method = "lm", trControl = train.control)
```

```
# 5 Point Summary for resulting RMSE for each resample  
summary(linearModel2$resample$RMSE)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
##    7.400   8.802   9.196   9.245   9.619  11.801
```

```
# 5 Point Summary for resulting R-Squared for each resample
summary(linearModel2$resample$Rsquared)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.5322  0.6239  0.6431  0.6425  0.6617  0.7235
```

```
# Histogram of RMSE values
```

```
hist(linearModel2$resample$RMSE, xlab = "RMSE Values", main = "Histogram of the RMSE values")
```



```
# Calculating the MSE from RMSE
mse2 = mean(linearModel2$resample$RMSE)^2

# Summarize the results
cat("Training Mean MSE (Bootstrap): ", mse2)
```

```
## Training Mean MSE (Bootstrap): 85.46133
```

```
cat("Training Mean RMSE (Bootstrap): ", mean(linearModel2$resample$RMSE))
```

```
## Training Mean RMSE (Bootstrap): 9.24453
```

```
cat("Training Mean R-squared (Bootstrap): ",mean(linearModel2$resample$Rsquared))
```

```
## Training Mean R-squared (Bootstrap): 0.6424873
```

```
predVals_boot = predict(linearModel2,testData)
mse3 = MSE(testData$Residuary, predVals_boot)
```

```
RSS = function (yActual, yPred)
{
  return (sum((yActual - yPred)^2))
}

TSS = function (yActual)
{
  return (sum((yActual - mean(yActual))^2))
}
```

```
rss = RSS(testData$Residuary, predVals_boot)
tss = TSS(testData$Residuary)

# Summarize the results
cat("Testing MSE (Bootstrap): ", mse3)
```

```
## Testing MSE (Bootstrap): 80.36479
```

```
cat("Testing RMSE (Bootstrap): ", sqrt(mse3))
```

```
## Testing RMSE (Bootstrap): 8.964641
```

```
cat("Testing Mean R-squared (Bootstrap): ",1 - (rss/tss))
```

```
## Testing Mean R-squared (Bootstrap): 0.6435649
```

On the test set, there is no difference in performance between the original and bootstrap models.

#Problem 4

```
library(readr)
library(data.table)
library(caret)
```

```
gcd_URL = "https://archive.ics.uci.edu/ml/machine-learning-databases/statlog/german/german.data-numeric"
gc_data = fread(gcd_URL, header = FALSE)
```

```
#As the response variable need to be numeric, converting categorical V25 to a factor  
gc_data$V25 = factor(gc_data$V25)
```

```
#Creating 80-20 Training Testing Split, createDataPartition() returns the indices  
trainIndex = createDataPartition(y = gc_data$V25 , p = 0.8, list = FALSE)
```

```
#Training data  
trainData = gc_data[trainIndex,]  
  
#Testing data (note the minus sign)  
testData = gc_data[-trainIndex,]
```

```
# Creating model for y = V25 using glm  
logisticModel1 = glm(V25~.,family=binomial,data=trainData)
```

```
actualVals = trainData$V25  
  
# Using a 50% cut-off factor i.e probabilities > 0.5 are 2 and rest are 1  
fittedVals = ifelse(logisticModel1$fitted.values > 0.5,2,1)  
fittedVals = factor(fittedVals)
```

```
# Confusion matrix  
cm = confusionMatrix(fittedVals, trainData$V25)  
cm
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   1   2
##           1 505 112
##           2  55 128
##
##           Accuracy : 0.7912
##           95% CI : (0.7614, 0.8189)
##           No Information Rate : 0.7
##           P-Value [Acc > NIR] : 3.653e-09
##
##           Kappa : 0.4668
##
##           McNemar's Test P-Value : 1.468e-05
##
##           Sensitivity : 0.9018
##           Specificity : 0.5333
##           Pos Pred Value : 0.8185
##           Neg Pred Value : 0.6995
##           Prevalence : 0.7000
##           Detection Rate : 0.6312
##           Detection Prevalence : 0.7712
##           Balanced Accuracy : 0.7176
##
##           'Positive' Class : 1
##
```

```
# Summarize the results
cat("Training Precision: ", cm$byClass[5] * 100, "%")
```

```
## Training Precision: 81.84765 %
```

```
cat("Training Recall: ", cm$byClass[6] * 100, "%")
```

```
## Training Recall: 90.17857 %
```

```
cat("Training F1-Score: ", cm$byClass[7] * 100, "%")
```

```
## Training F1-Score: 85.81138 %
```

```
# Predict [By setting the parameter type='response', R will output probabilities in the form of
P(y=1|X)]
probs = predict(logisticModel1, testData, type = "response")

#Using a 50% cut-off factor i.e probabilities > 0.5 are Males and rest are Females
fittedVals_test = ifelse(probs > 0.5,2,1)
fittedVals_test = factor(fittedVals_test)

# Confusion matrix
cm_test = confusionMatrix(fittedVals_test, testData$V25)
cm_test
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    1    2
##           1 122  29
##           2  18  31
##
##           Accuracy : 0.765
##           95% CI : (0.7, 0.8219)
##       No Information Rate : 0.7
##       P-Value [Acc > NIR] : 0.02493
##
##           Kappa : 0.4095
##
##  McNemar's Test P-Value : 0.14466
##
##           Sensitivity : 0.8714
##           Specificity : 0.5167
##       Pos Pred Value : 0.8079
##       Neg Pred Value : 0.6327
##           Prevalence : 0.7000
##       Detection Rate : 0.6100
##  Detection Prevalence : 0.7550
##       Balanced Accuracy : 0.6940
##
##       'Positive' Class : 1
##
```

```
# Summarize the results
cat("Testing Precision: ", cm_test$byClass[5] * 100, "%")
```

```
## Testing Precision:  80.7947 %
```

```
cat("Testing Recall: ", cm_test$byClass[6] * 100, "%")
```

```
## Testing Recall:  87.14286 %
```

```
cat("Testing F1-Score: ", cm_test$byClass[7] * 100, "%")
```

```
## Testing F1-Score: 83.8488 %
```

#—————Cross-Validation—————

```
# Define training control
```

```
train.control = trainControl(method = "cv", number = 10)
```

```
# Train the model
```

```
logisticModel2 = train(V25~., data = trainData, method = "glm", family = "binomial", trControl =  
train.control)
```

```
fittedVals_cv = ifelse(logisticModel2$finalModel$fitted.values > 0.5,2,1)
```

```
fittedVals_cv = factor(fittedVals_cv)
```

```
# Confusion matrix
```

```
cm_cv = confusionMatrix(fittedVals_cv, trainData$V25)
```

```
cm_cv
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  1    2
```

```
##           1 505 112
```

```
##           2   55 128
```

```
##
```

```
##           Accuracy : 0.7912
```

```
##           95% CI : (0.7614, 0.8189)
```

```
## No Information Rate : 0.7
```

```
## P-Value [Acc > NIR] : 3.653e-09
```

```
##
```

```
##           Kappa : 0.4668
```

```
##
```

```
## McNemar's Test P-Value : 1.468e-05
```

```
##
```

```
##           Sensitivity : 0.9018
```

```
##           Specificity : 0.5333
```

```
## Pos Pred Value : 0.8185
```

```
## Neg Pred Value : 0.6995
```

```
## Prevalence : 0.7000
```

```
## Detection Rate : 0.6312
```

```
## Detection Prevalence : 0.7712
```

```
## Balanced Accuracy : 0.7176
```

```
##
```

```
## 'Positive' Class : 1
```

```
##
```

```
# Summarize the results
```

```
cat("Training Precision with 10-fold CV: ", cm_cv$byClass[5] * 100, "%")
```



```
## Training Precision with 10-fold CV: 81.84765 %
```

```
cat("Training Recall with 10-fold CV: ", cm_cv$byClass[6] * 100, "%")
```

```
## Training Recall with 10-fold CV: 90.17857 %
```

```
cat("Training F1-Score with 10-fold CV: ", cm_cv$byClass[7] * 100, "%")
```

```
## Training F1-Score with 10-fold CV: 85.81138 %
```

```
# Predict [By setting the parameter type='response', R will output probabilities in the form of  
P(y=1|X)]
```

```
probs_cv = predict(logisticModel2, testData, type = "prob")
```

```
#Using a 50% cut-off factor i.e probabilities > 0.5 are Males and rest are Females
```

```
fittedVals_cv_test = ifelse(probs > 0.5,2,1)
```

```
fittedVals_cv_test = factor(fittedVals_test)
```

```
# Confusion matrix
```

```
cm_cv_test = confusionMatrix(fittedVals_test, testData$V25)
```

```
cm_cv_test
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   1   2
##           1 122  29
##           2  18  31
##
##           Accuracy : 0.765
##           95% CI : (0.7, 0.8219)
##       No Information Rate : 0.7
##       P-Value [Acc > NIR] : 0.02493
##
##           Kappa : 0.4095
##
##  McNemar's Test P-Value : 0.14466
##
##           Sensitivity : 0.8714
##           Specificity : 0.5167
##       Pos Pred Value : 0.8079
##       Neg Pred Value : 0.6327
##           Prevalence : 0.7000
##       Detection Rate : 0.6100
##   Detection Prevalence : 0.7550
##       Balanced Accuracy : 0.6940
##
##       'Positive' Class : 1
##
```

```
# Summarize the results
cat("Testing Precision: ", cm_cv_test$byClass[5] * 100, "%")
```

```
## Testing Precision:  80.7947 %
```

```
cat("Testing Recall: ", cm_cv_test$byClass[6] * 100, "%")
```

```
## Testing Recall:  87.14286 %
```

```
cat("Testing F1-Score: ", cm_cv_test$byClass[7] * 100, "%")
```

```
## Testing F1-Score:  83.8488 %
```

On the test set, there is no difference in performance between the original and bootstrap models.